

DWIJ PRAJAPATI

AI & Backend Engineer | dwijprajapati46476@gmail.com | +91-9979246476

GitHub: github.com/dwijprajapati | LinkedIn: linkedin.com/in/dwij-prajapati

SUMMARY

AI & Backend Engineer building real-time voice AI agents, LLM-as-a-Judge evaluation pipelines, and high-performance ML platforms with sub-second latency and resilient stateful architecture.

EXPERIENCE

Banao Technologies (InterviewGod.ai) AI Developer Intern (Feb 2026 - Present)

- Architected event-driven voice pipeline on LiveKit (WebRTC), GPT-4o, ElevenLabs and Sarvam AI (<1s latency).
- Designed adaptive multi-turn dialog state machine with repeat-request & filler-word interceptors.
- Built LLM-as-a-Judge validity gate (gpt-4o-mini) with fail-open resiliency to eliminate hallucinations.
- Engineered Redis-backed session checkpointing with 24-hour buffer for zero-data-loss reconnects.
- Deployed multi-service Docker Compose stacks to AWS EC2 with Caddy TLS and CI/CD pipelines.

PROJECTS

FinSight AI Quantitative Intelligence & Trading Platform

- Combined SARIMAX, regularized LSTM/MLP networks, and FinBERT sentiment scoring.
- Built MDP trading environment training PPO, A2C, and DQN reinforcement learning agents.
- Generated Efficient Frontier via SLSQP Markowitz optimization and 10,000+ Monte Carlo runs.

CompInsight AI Enterprise ML Compensation Modeling Engine

- Benchmarked CatBoost, LightGBM, and XGBoost with Optuna Bayesian hyperparameter search.
- Achieved 0.93 R-squared score with 22% reduction in RMSE over baseline regression.
- Exposed feature attribution with SHAP TreeExplainer and served sub-40ms P99 inference via FastAPI.

TECHNICAL SKILLS

Languages: Python, C++, C, JavaScript, TypeScript, SQL, Bash/Shell

AI & ML: LLM Orchestration, Prompt Scaffolding, LiveKit WebRTC, STT/TTS, Scikit-learn, TensorFlow

Backend & Cloud: FastAPI, Django, Asyncio, Redis, PostgreSQL, Docker, AWS (EC2, S3), Git, Pytest

EDUCATION

B.Tech in Computer Engineering G.H. Patel College of Engineering (2022 - 2026) | CGPA 8.18 / 10